

# Towards a new framework for integration of network planes


**Siddhant Ray**

Semester Thesis Presentation

25th June 2021, Zürich



# Traditional Layer-3 router:



# Programmable data planes : Why should we care?

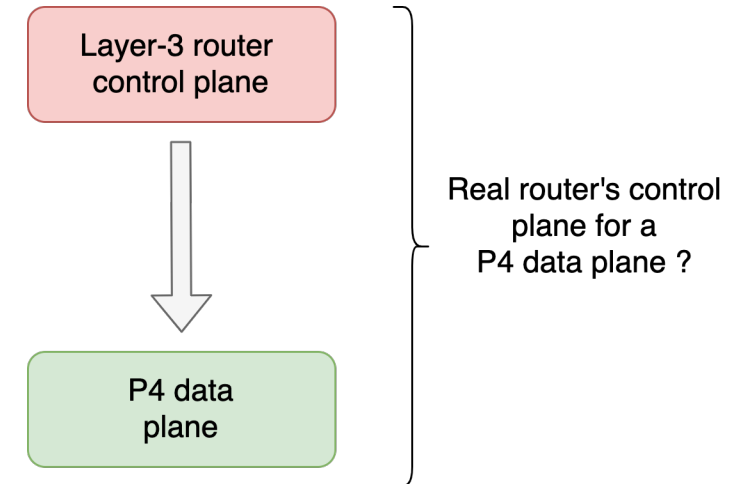
- Data planes have usually been **static** for eg. \*NIX kernels
- The control plane **tells** the data plane **how to forward packets**
- With programmable data planes , we can
  - Make intelligent forwarding decisions at the data plane level
  - Greater packet level control over forwarding
  - Send feedback to control plane, custom forwarding etc.

# Programmable data planes : Challenges

- Existing P4 data planes : do not have supporting routing protocols
- Need a custom controller : compute routes, populate data plane
- Limited : Doesn't provide all properties of actual routing

So, can we do better?

Can we integrate existing real routing protocols to work with P4 data planes?





# In this project, we present :

- A new **Super-Node** for Layer-3 forwarding
- Each Super-Node has :
  - An FRR control plane
  - A P4 programmable data plane
- We create and run :
  - FRR control planes in **individual Linux namespaces**
  - P4 data planes in a **common root namespace**
  - Protocols **OSPF and iBGP** for Super-Nodes in the same Autonomous System (AS), **eBGP** on Super-Nodes at the border of an AS.
  - Load balancing for end-to-end forwarding

# Framework: Mininet [1] and p4-utils [2]

## ➤ Mininet :

- A rapid, prototyping tool which helps create emulated networks inside a single OS kernel
- Uses the Linux kernel and network stack, nodes can be in the kernel or custom network namespaces.
- Our control plane routers run as custom Mininet nodes.

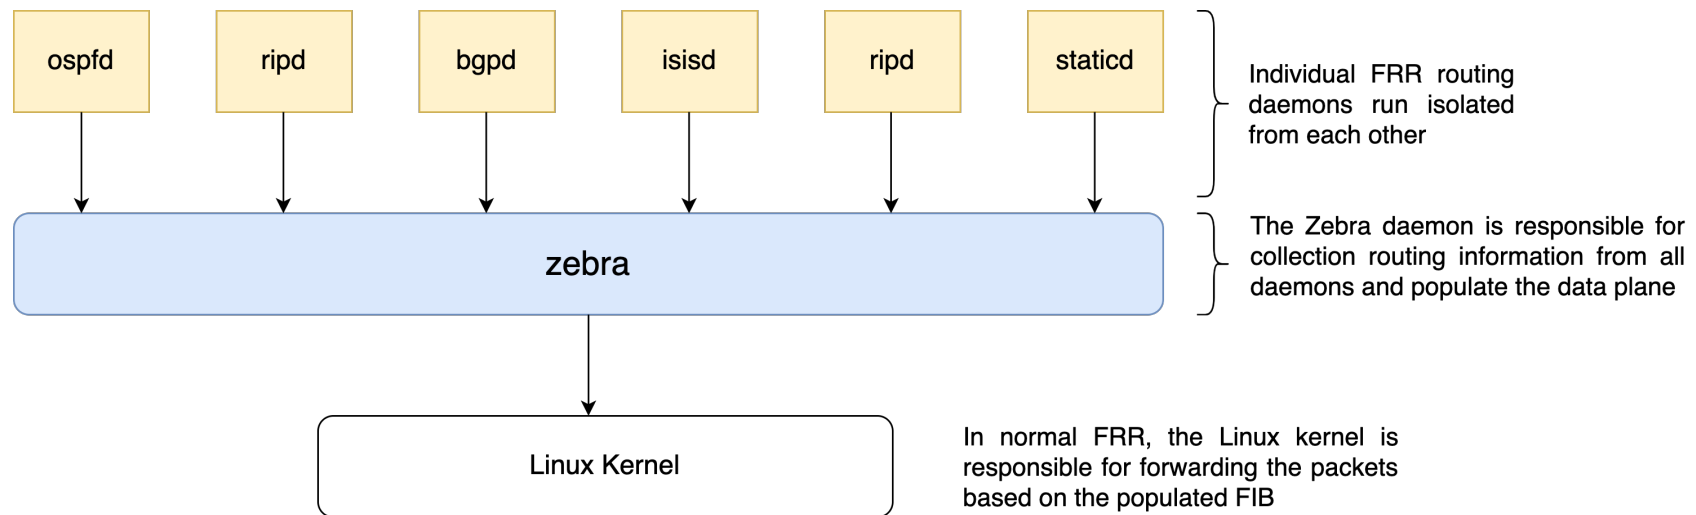
## ➤ p4-utils :

- Used to build and develop P4 networks. It is an extension to Mininet.
- The P4 data plane is implemented as a custom Mininet switch.

[1] <http://mininet.org/>

[2] <https://github.com/nsg-ethz/p4-utils>

# Control Plane : FRRouter [3]



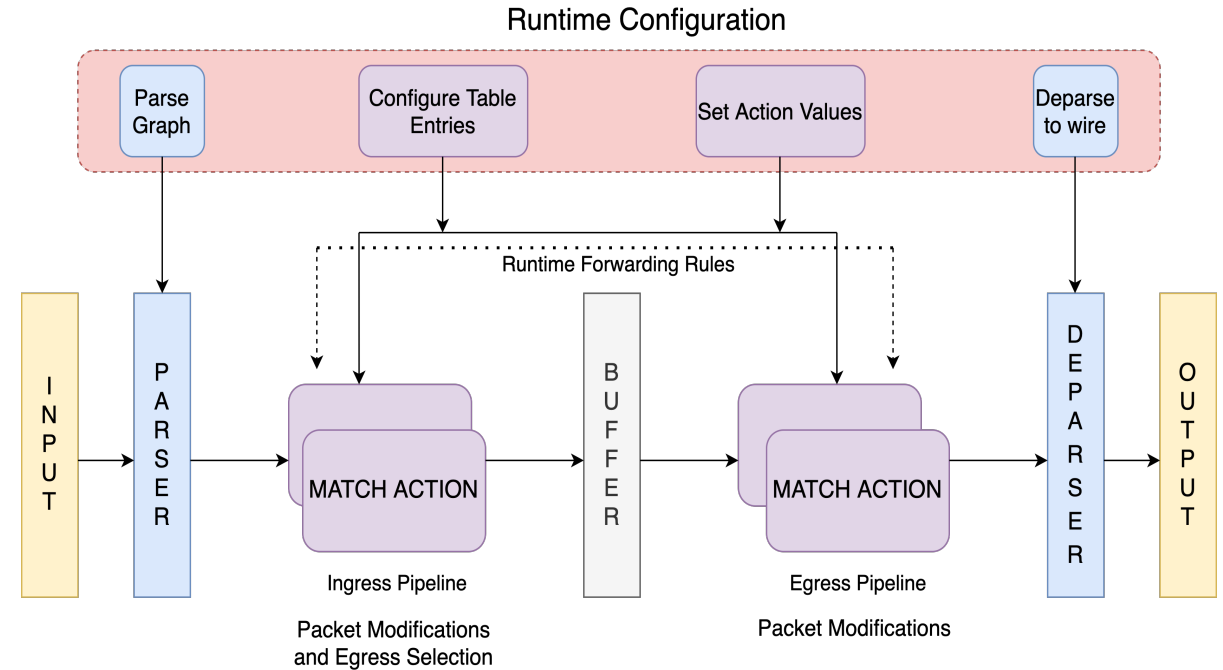
FRR suite architecture

[3] <https://frrouting.org/>

# Data plane : P4 bmv2 Switch [4]

A P4 data plane consists of a :

- Parser :  
Extract information from the packet header
- Match-Action :  
Match on fields to set actions to decide forwarding
- Deparser :  
Add parsed headers back at the end



P4 data plane pipeline

[4] <https://p4.org/index.html>




# Initial Super-Node design :

- FRR control plane
- P4 data plane

The control planes routers do not have **real links** to each other, they have non-connected **fake interfaces**.

The P4 data planes act as “**transparent forwarding**” nodes for the control planes.



# However there is a problem..

- FRR routers **do not send packets** on interfaces with **no real links**.
- FRR routers **drop packets** if not received on the **intended interface**.

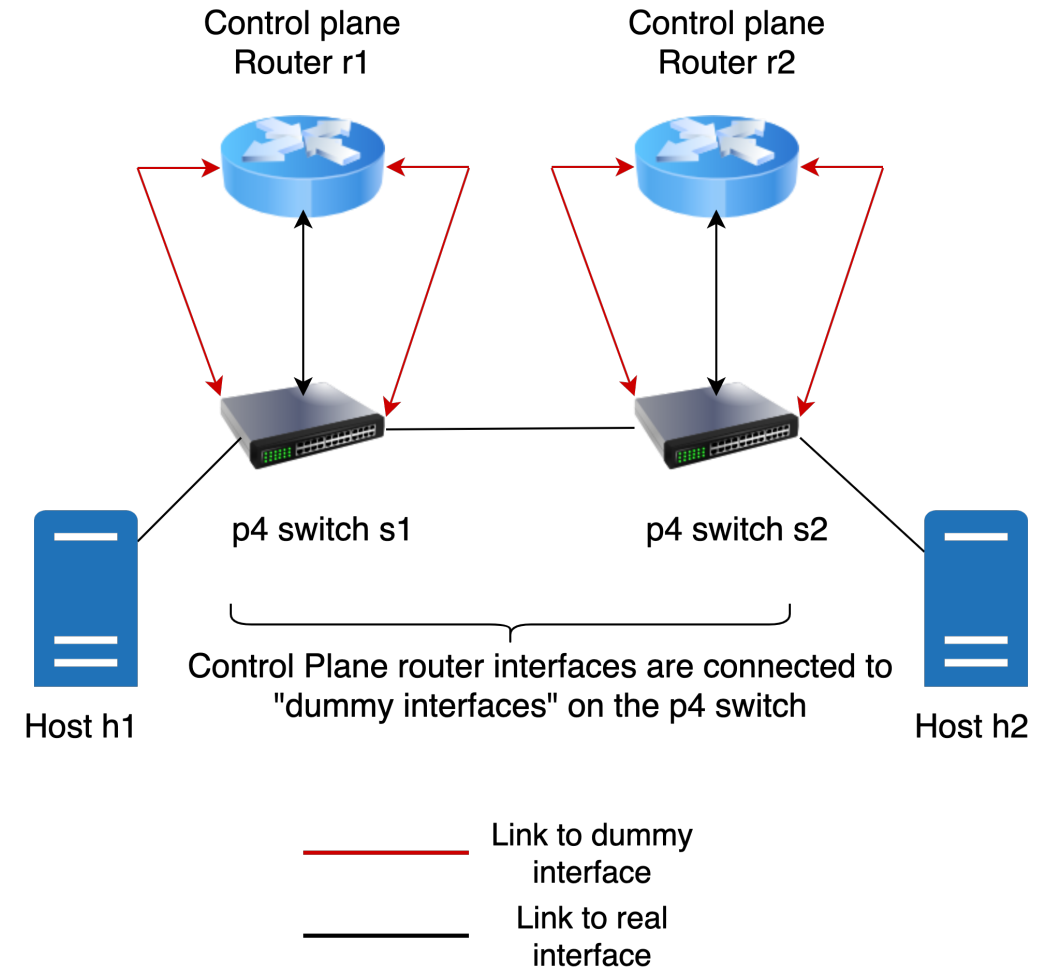
Two options :

- Add dummy links to nothing , capture packets and reinject to correct interface.  
(This approach is **very slow**)
- Add dummy links to the P4 data plane, no need to reinject packets.  
(**Fast** but more interfaces on the P4 data plane)

We choose the second option for our design.

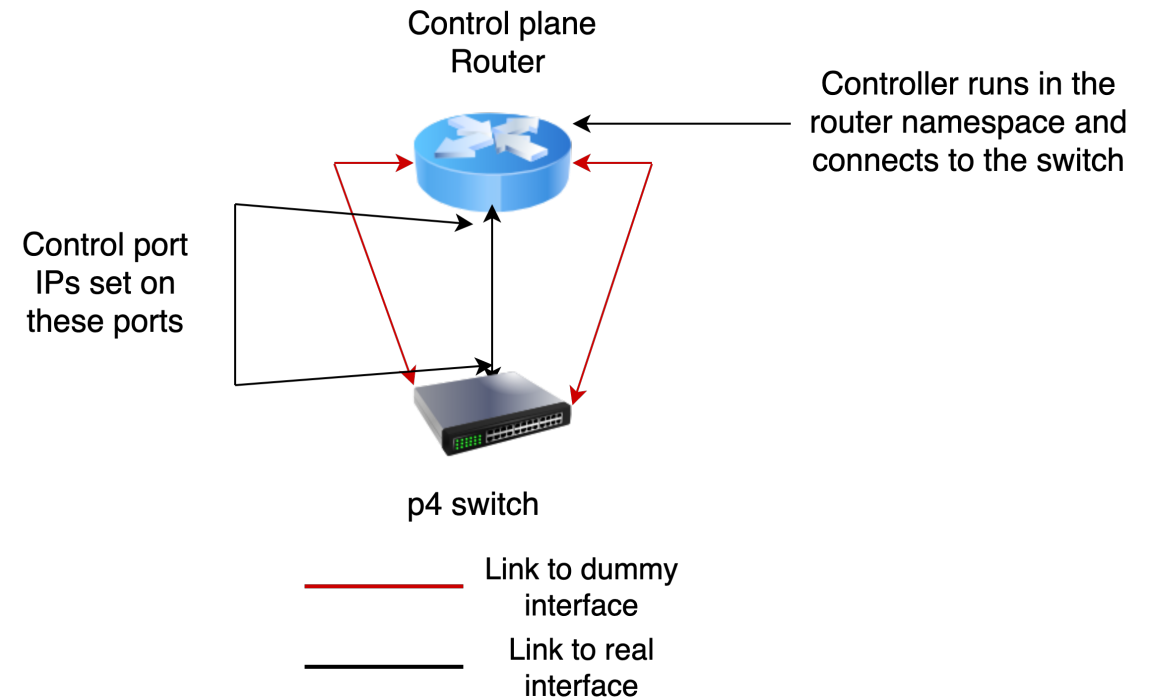
# Final Super-Node design :

- Control planes routers have dummy links to each P4 switch.
- The FRR control plane router is fooled to believe it has a real connection to the other.
- Control plane routers send and receive the correct packets.



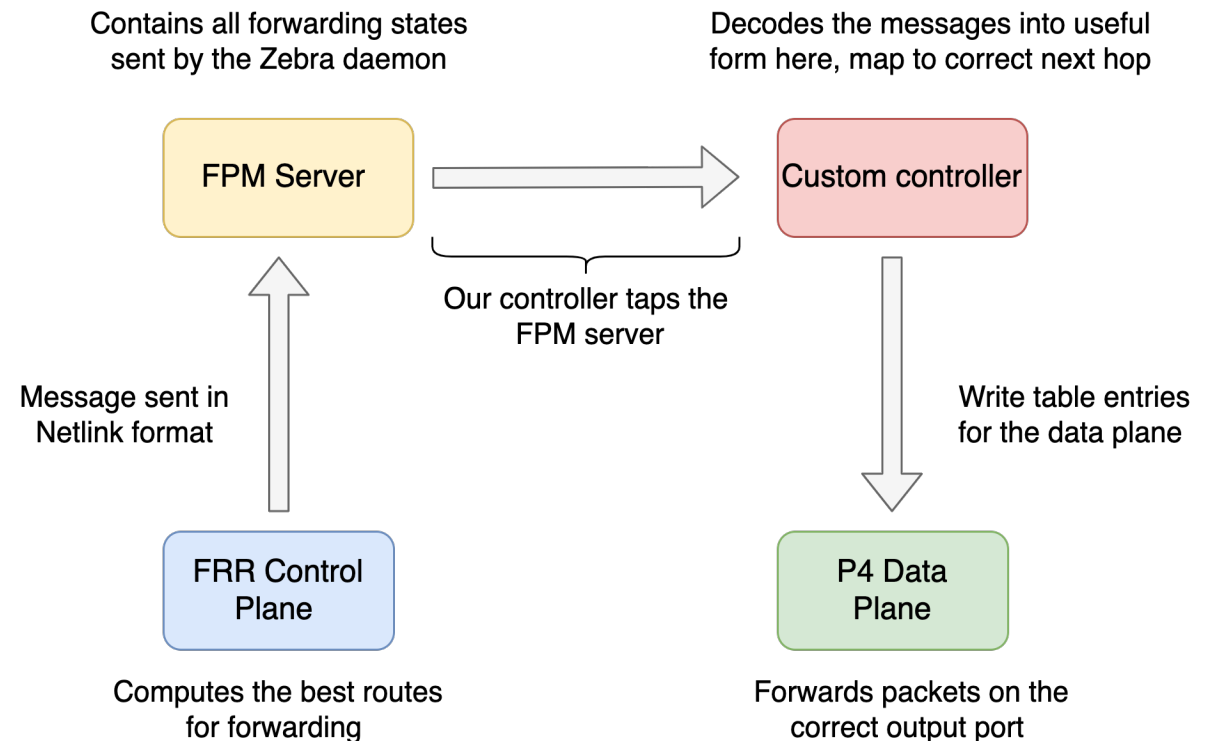
# Handling data plane updates :

- A Python controller running in each control plane router namespace.
- A P4 program running in the P4 data plane.
- The P4 data plane forwards control plane packets from dummy links to the correct output interface. (for OSPF, BGP etc.)
- Control plane routers now calculate best routes to all destinations.



# Handling data plane updates :

- FRR has an FPM server which collects a complete copy of the forwarding table.
- FPM by itself doesn't push entries into the data plane.
- Our controller is dynamic: updates routes into the data plane as soon as it arrives.
- Our programmable data plane ensures: only IPv4 packets are forwarded using FPM routes.



# Understanding FPM messages :

## Sample decoded FPM message :

```
{'family': 2, 'dst_len': 24, 'src_len': 0, 'tos': 0, 'table':  
254, 'proto':  
11, 'scope': 0, 'type': 1, 'flags': 0, 'attrs': [('RTA_DST',  
'10.2.0.0'),  
(('RTA_PRIORITY', 11), ('RTA_GATEWAY',  
'10.0.1.2'), ('RTA_OIF', 26715))],  
'header': {'length': 60, 'type': 24, 'flags': 1025,  
'sequence_number': 0,  
'pid': 0}}
```

Message to reach host h2 from Super-Node r1-s1 (no multipath configured).

## What our controller does:

Check family :

2 > Means IPv4 route

Check destination prefix :

24 > Inside AS, 8 > To other AS

Extract destination IP prefix :

10.2.0.0/24

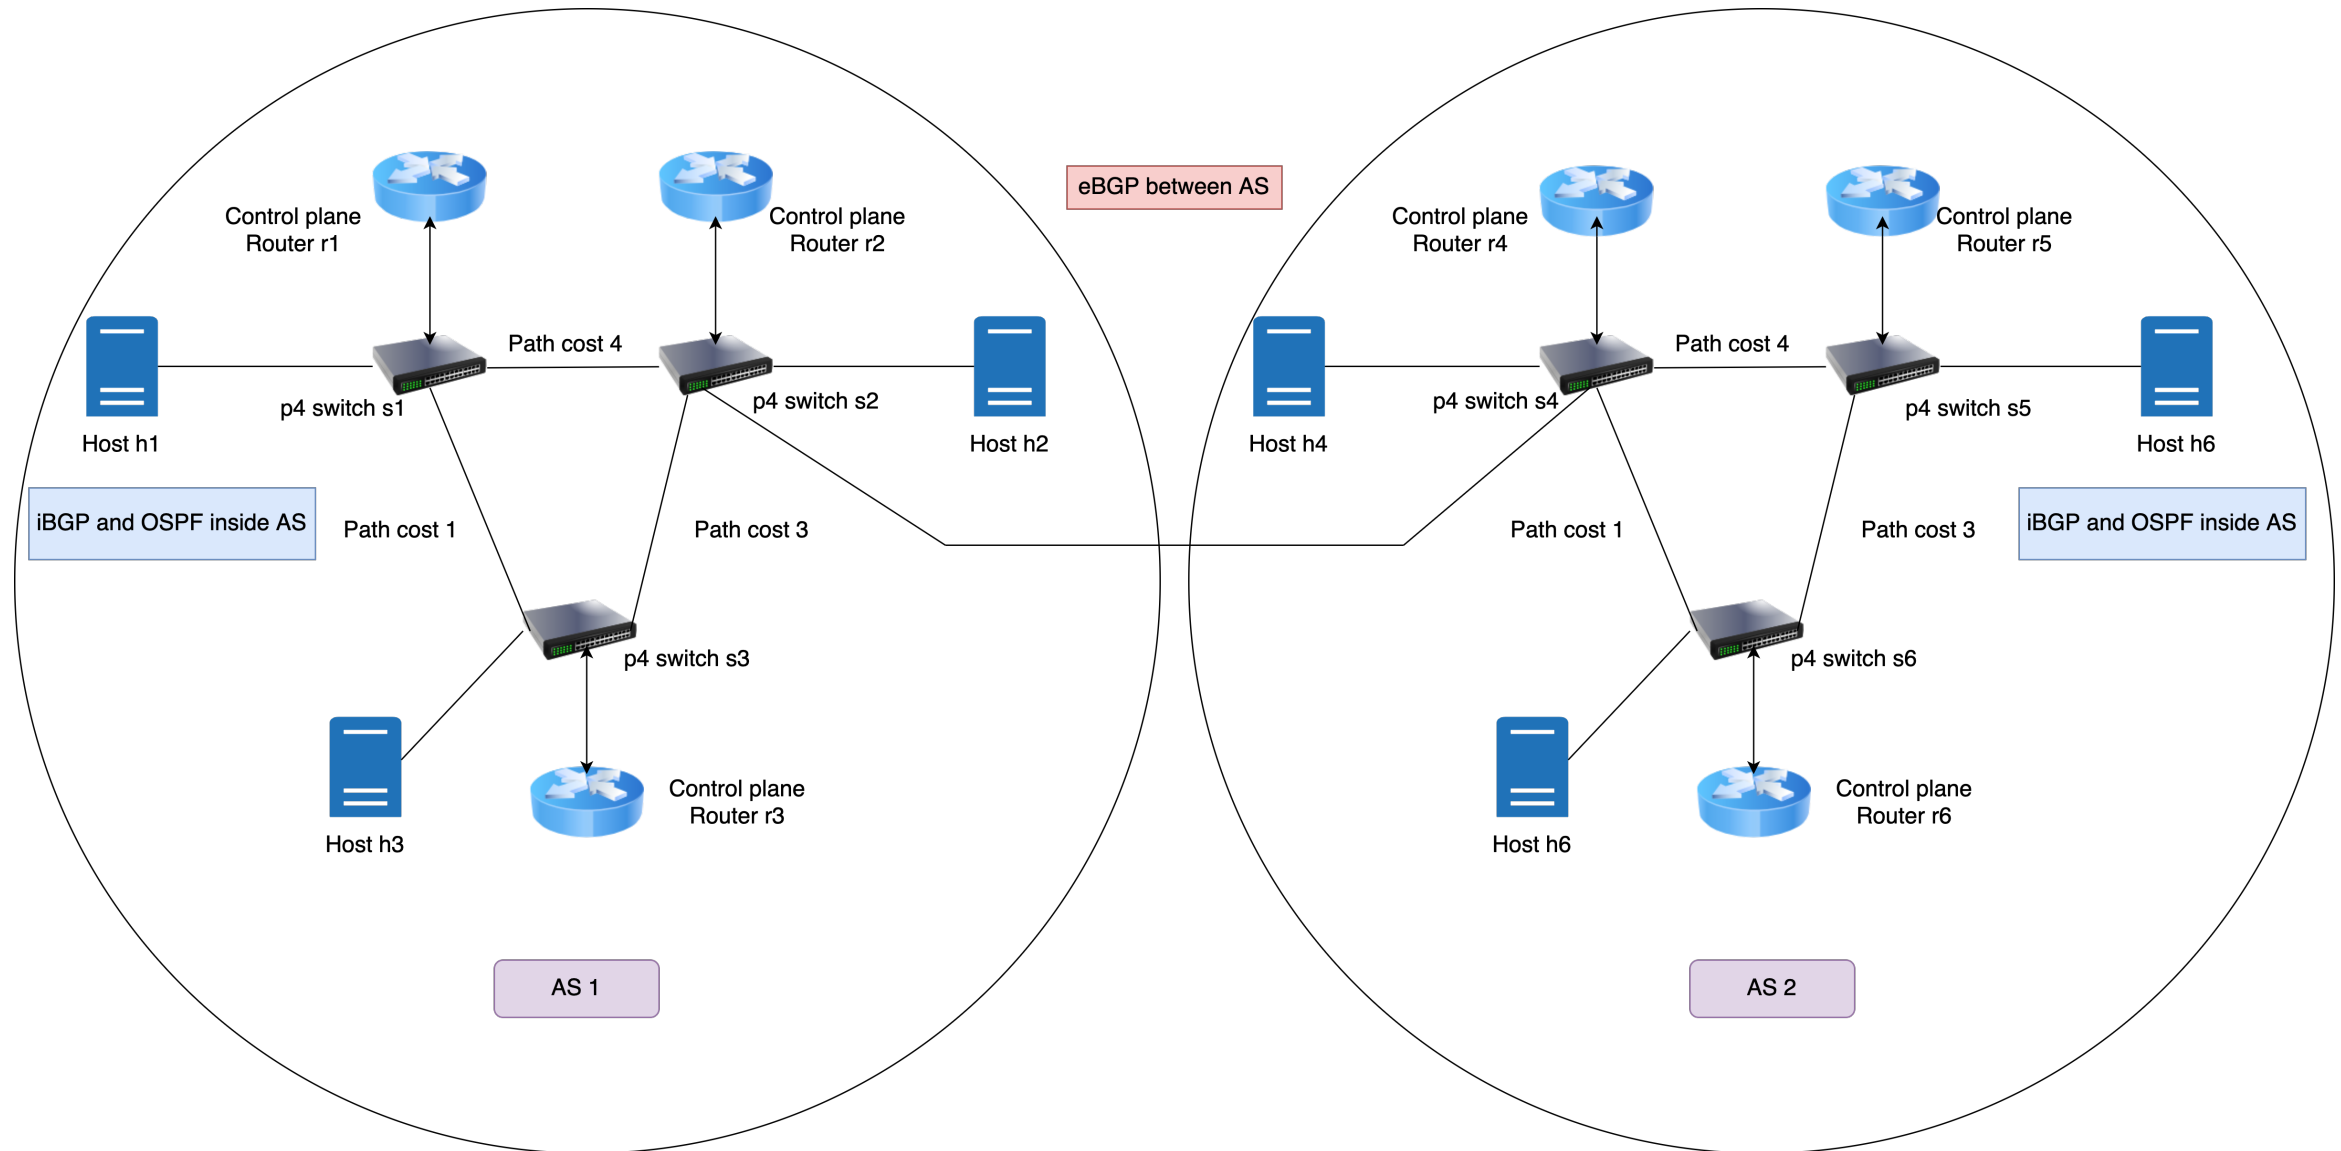
Map to correct output port :

26715 > 2

Write correct table entry for IPv4 forwarding

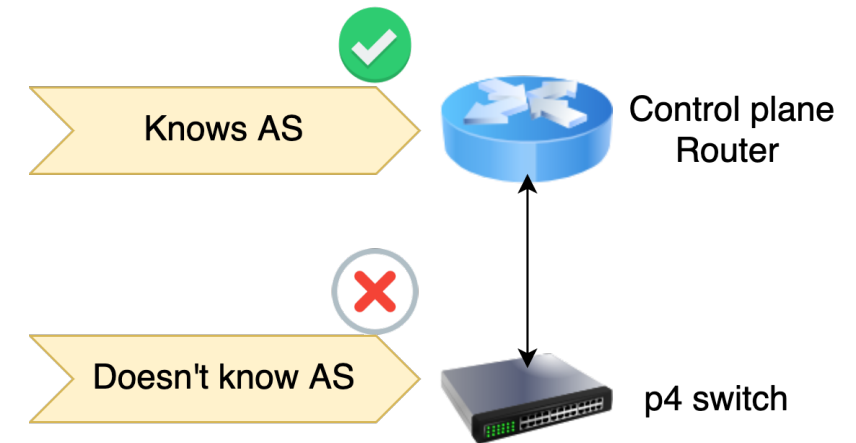


# Testing Topology :



# Feature : Running OSPF and BGP


- AS - only known to control plane routers, not to the P4 data plane
- Super-Nodes at the border : run eBGP
- Interior Super-Nodes : run iBGP and OSPF
- eBGP learnt routes are distributed over iBGP (preferred) or OSPF.
- Sustained TCP connection over the eBGP peers for connected BGP state
- OSPF – updates internal AS routes
- BGP – updates external AS routes



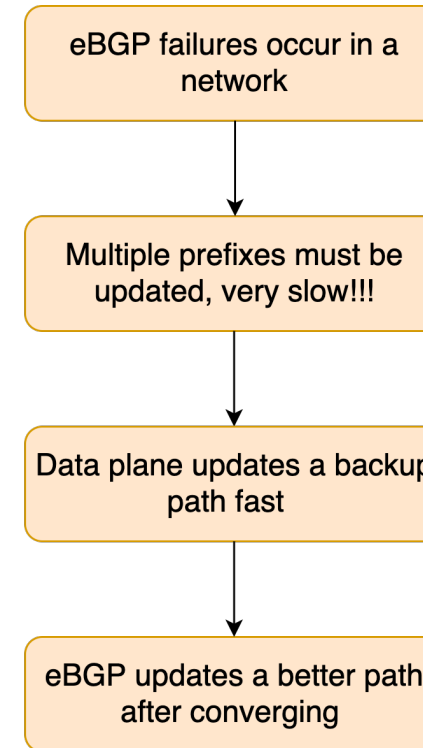
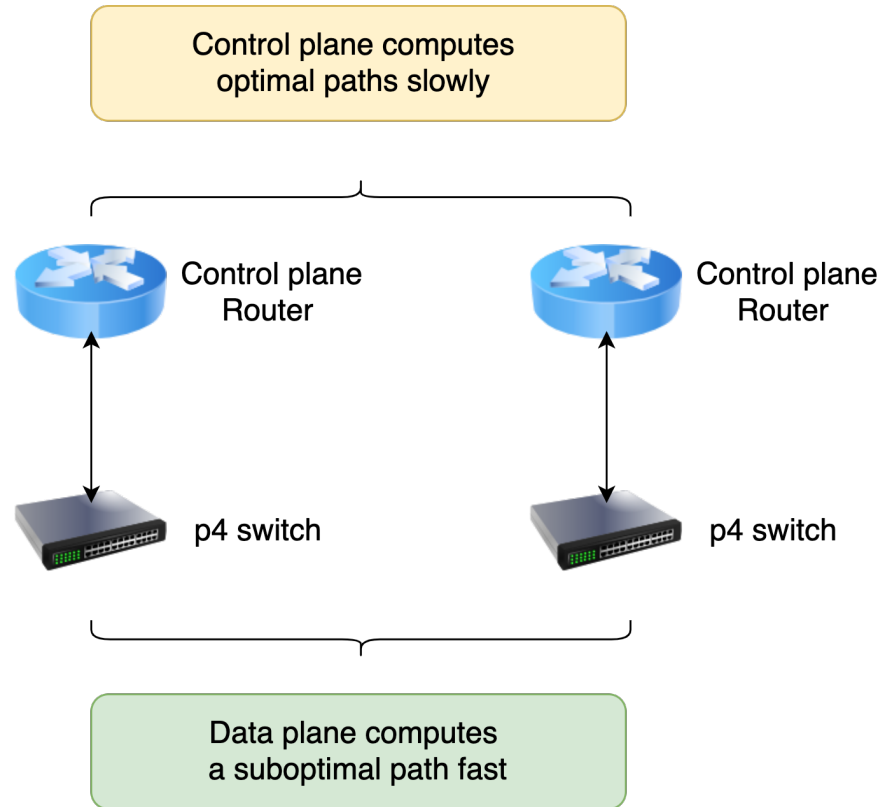
Super-Node setup

# Feature : Load Balancing using ECMP

- FRR identifies multipath routing on equi-cost paths.
- P4 program in the data plane calculates ECMP hash values on the standard 5-tuple.
- Controller checks for multi-path and populates the ECMP table for multi-path forwarding.
- Different ECMP groups for different IP destination prefixes (per AS) in order to have multi-path across AS.



# What can we do now?



# Outlook :

- Limited to adding dummy links externally via Linux :  
Currently no support for multiple, parallel links via Mininet or p4-utils
- Need to rely on dummy links :  
FRR doesn't send packets on non-connected interfaces
- Cannot stop FRR from populating the kernel data plane  
P4 data planes can perform “transparent forwarding” on control plane packets
- Control plane must forward a lot of packets via the P4 data plane  
Programmability ensures complete control to have correct forwarding



Time for questions



Siddhant Ray  
MSc. in Electrical Engineering and Information Technology  
[sidray@student.ethz.ch](mailto:sidray@student.ethz.ch)

D-ITET  
ETH Zurich